

ETEC3702 – Concurrency

Final Exam Study Guide

Topics

You are responsible for **everything** covered in class through Thursday, 23 April 2020.

Some of the major (but not necessarily all) topics included will be:

- Introduction to Concurrency
 - Definition
 - Concurrent Versus Sequential
 - Formal notation
 - “For all“
 - “There exists”
 - “It is true that”
 - “Concurrent with” (||)
 - “Precedes” (→)
 - Formal definition of sequential
 - Formal definition of concurrent
 - Non-deterministic execution
- Dangers of Concurrency
 - Effective orderings – how to compute
 - Problematic orderings
 - Safety Property Violations
 - Concurrent update problem (turnstiles problem, bank account problem)
 - Critical Sections
 - Mutual Exclusion
 - Software Solutions
 - Dekker’s Algorithm
 - Peterson’s Algorithm
 - Liveness Properties Violations
 - Deadlock
 - Livelock
 - Starvation

- Difficulty in Testing
- Motivation for Concurrent Solutions
 - Organization
 - Natural Solutions
 - Relax over-specification
 - Multiple separate tasks
 - Speed-up
 - Moore's Law and the Historic Trend of CPU Clock Speeds
 - Multiple Core CPUs
 - Distributed Computing / "Cloud" Computing
 - Specialty Parallel Processors / GPU Computing
 - Computing Speedup to compare sequential versus concurrent solutions.
 - Computing Maximum speedup of a program with sequential and parallelizable parts.
 - Unavoidable / Inherent
 - Network / Communication Systems
 - Database Systems
 - Distributed Systems
 - Event-driven Systems
 - Physical Systems.
- Synchronization mechanisms
 - Locks
 - acquire()
 - release ()
 - with lock syntax.
 - Optional arguments (blocking, timeout)
 - The "Monitor" object design pattern concept.
 - Rlocks
 - motivation for / versus normal locks.
 - acquire()
 - release()
 - with rlock syntax
 - Optional arguments (blocking, timeout)
 - Semaphores

- As a data-type
- Set of values
- Set of permissible operations
 - Definition of P() or acquire()
 - Definition of V() or release()
- Types of Semaphores
 - Binary
 - Counting
- In Python
 - Semaphore objects
 - BoundedSemaphore objects

□ Conditional Sections

- Producer / Consumer design patterns.
- The need for conditional synchronization
- Condition objects
 - acquire()
 - release()
 - wait()
 - notify()
 - notify_all()
 - Use to implement a “Monitor” design pattern

□ Events

- Concept and function
- As a synchronization mechanism
- As a communication mechanism
- Event objects
 - set()
 - clear()
 - is_set()
 - wait()

□ Barriers

- Concept and function
- Uses
- Barrier objects
 - wait()

- abort()
- reset()
- broken
- parties
- n_waiting

○ Communication Mechanisms

□ Global variables

□ Events

□ Queues

- Queue(), LifoQueue(), PriorityQueue()
- Optional args: maxsize
- put()
- get()
- empty()
- full()
- task_done()
- join()

○ Multiprocessing

□ Differences from threading

□ Advantages over threading

□ Disadvantages related to threading

□ Potential for speedup

□ Need for the `__name__ == "__main__"` check

□ `cpu_count()`

□ Inter-process communication using queues.

□ Using queues with processes to “return” values.

□ Support for other synchronization and communication mechanisms

- Lock
- Rlock
- Condition
- Semaphore / BoundedSemaphore
- BarriersEvent
- Queue

○ Inter-process Communication

❑ Queues

❑ Pipes

- Use in operating systems to allow IPC
- multiprocessing.Pipe()
- send()
- recv()
- close()
- pipes vs. queues

❑ Managers

- multiprocessing.Manager()
- manager namespaces
- adding items to a manager namespace
- immutable versus mutable objects in managers
- manager.list() type

❑ Value / Array Objects

- multiprocessing.Value() and typecodes
- multiprocessing.Array()
- The need to lock while accessing Value() / Array() data.
- Implementing a Monitor to protect Value / Array objects.

❑ Shared memory

- multiprocessing.shared_memory
- creating shared_memory, create, size, name
- advantages and disadvantages wrt other IPC mechanisms.
- Connecting a process to a named shared_memory block.
- Using struct.pack and struct.unpack to store data into shared_memory.

❑ IPC for distributed computing

- Sockets
 - creating a socket object
 - connect()
 - bind()
 - listen()
 - accept()
 - send() / recv()
 - close()

- Encoding / decoding socket data
 - binary / custom
 - JSON
 - Pickle
 - XML
 - MsgPack
 - encoding / decoding with loads() / dumps()
- multiprocessing.connection objects
 - advantages over sockets
 - send() / recv() / poll()
 - send_bytes() / recv_bytes()
 - Listener() / Client()
 - sending objects over connection objects

□ Event-driven Programming and asyncio

- cooperative multi-tasking.
- concept and use with concurrent systems
- asyncio
- event loop concepts and use with I/O-bound tasks
- async / await syntax and uses.
- Awaitable items: coroutines, tasks, futures.
- async def coroutine functions.
- Tasks and asyncio.run()
- asyncio.wait()
- asyncio.sleep()
- asyncio.gather()
- futures

□ Liveness: Deadlock and Livelock

- Liveness definition / Safety Definition
- Deadlock
 - Deadlock definition
 - 4 required conditions
 - Deadlock Prevention
 - Deadlock Detection and Recovery
 - Resource Allocation Graphs
 - Reduction to detect deadlock

- recovery methods
- Deadlock avoidance
 - Safe vs. unsafe vs. deadlock states
 - Djikstra's Banker's Algorithm
- Livelock
 - livelock definition
 - differences from deadlock
- The dining philosophers example
 - deadlock
 - livelock

□ Concurrency in C and C++

- Posix Threads / pthreads
 - pthread.h
 - pthread_create()
 - pthread_join()
 - pthread_exit()
 - pthread synchronization
 - mutexes
 - pthread_mutex_lock()
 - pthread_mutex_unlock()
 - condition variables
 - pthread_cond_wait()
 - pthread_cond_signal()
 - pthread_cond_broadcast()
- C++11 Concurrency
 - #include <thread>
 - creating thread object instances: thread t1()
 - .join()
 - thread synchronization
 - std::mutex
 - .lock()
 - .unlock()
 - std::condition_variable
 - Other mechanisms

Reminders/Advice:

- The entire exam will be open-notes.
- The exam will be posted on Blackboard on the day of the exam.
- The exam will have a time-limit and must be completed and submitted within that time limit.
- There may be a programming portion of the exam. You can use a computer for that portion, but you may not communicate with others during that portion. (no discord, no stack-exchange, no reddit, no facebook messenger, no zoom, no SMS messages, no whatsapp, no email, no message boards, no talking, no communicating whatsoever!) Seriously, if I find evidence that you've talked to others or been intellectually dishonest in any way you will fail the exam.
- You should complete and understand all of the assigned labs.
- You may have to write and submit some code. Be prepared to do that.
- Open notes, but don't assume open-book == no preparation/study.